

# GPT COST

Updated on 8 May 2026

Source: <https://docs.google.com/document/d/1FybSxB3lXJ5NOyA-EVsAm-Cz9Wpv0Yydrmsn8SeefRo/edit?usp=sharing>

|           |            |                               |              |                |
|-----------|------------|-------------------------------|--------------|----------------|
| #Messages | 50,000,000 | RATE (3 responses, in use)    | \$368,072.50 | ¥59,771,477.31 |
|           |            | RATE (1 response, Hypothesis) | \$180,312.50 | ¥29,281,037.03 |

Inc. Voice input (Google STT)

|                |
|----------------|
| ¥73,580,193.30 |
| ¥43,089,753.02 |

| Model Rates (per 1M tokens) |        |
|-----------------------------|--------|
| gpt-4o-mini input           | \$0.15 |
| gpt-4o-mini output          | \$0.60 |
| gpt-4.1-mini input          | \$0.40 |
| gpt-4.1-mini output         | \$1.60 |
| gemini-2.5-flash input      | \$0.30 |
| gemini-2.5-flash output     | \$2.50 |

| Step                   | Model                      | Avg input tokens | Avg input cost | Avg output tokens | Avg output cost |
|------------------------|----------------------------|------------------|----------------|-------------------|-----------------|
| Summarize Conversation | gpt-4o-mini                | 21,850,000,000   | \$3,277.50     | 1,300,000,000     | \$780.00        |
| Execute Workflow       | gpt-4.1-mini (3 responses) | 612,500,000,000  | \$245,000.00   | 22,900,000,000    | \$36,640.00     |
|                        | gpt-4.1-mini (1 response)  | 204,166,666,667  | \$81,666.67    | 7,633,333,333     | \$12,213.33     |
|                        | gpt-4o-mini                | 178,000,000,000  | \$26,700.00    | 2,750,000,000     | \$1,650.00      |
|                        | gemini-2.5-flash           | 5,000,000,000    | \$53,400.00    | 250,000,000       | \$625.00        |

| EXECUTE WORKFLOW: |                       |                  |                     |
|-------------------|-----------------------|------------------|---------------------|
| #                 | Tasks                 | Model            | Notes               |
| 1                 | Language Detection    | gemini-2.5-flash |                     |
| 2                 | Query Rewrite         | gpt-4o-mini      |                     |
| 3                 | Parallel Response     | gpt-4.1-mini     | (parallel response) |
| 4                 | Generation            | gpt-4o-mini      |                     |
| 5                 | Evaluation (optional) | gpt-4o-mini      |                     |

## [OPTION] VOICE INPUT

| S2T Model Rates (per minutes) |         |
|-------------------------------|---------|
| GPT-Realtime-Whisper          | \$0.017 |
| Google Speech-to-Text         | \$0.016 |
| Amazon Transcribe             | \$0.024 |

| Estimated cost |             | Calculation Logic                                                        |
|----------------|-------------|--------------------------------------------------------------------------|
| \$90,348.64    | ¥14,671,761 | #message * (avg_input_syllables / avg_syllables_speed) * (Rate/ 60 secs) |
| \$85,034.01    | ¥13,808,716 |                                                                          |
| \$127,551.02   | ¥20,713,074 |                                                                          |

Cell: G32

Comment: Khanh Phan

avg\_input\_syllables: 50 syllables per input message.

avg\_syllables\_speed: 7.84 syllables per second (Japanese native speakers)

5/13/2026 12:19 AM